

## 1.3 BPF\_PROG\_TYPE\_SK\_SKB

- What do I do with it? Allows users to access skb and socket details such as port and IP address with a view to supporting redirect of skbs between sockets. See <https://lwn.net/Articles/731133/>. This functionality is used in conjunction with a sockmap – a special-purpose BPF map that contains references to socket structures and associated values. sockmaps are used to support redirection. The program is attached and the `bpf_sk_redirect_map()` helper can be used to carry out the redirection. The general approach we catch socket creation events with `sock_ops` BPF programs, associate values with the sockmap for these, and then use data at the `sk_skb` instrumentation points to inform socket redirection – this is termed the verdict, and the program for this is attached to the sockmap via `BPF_SK_SKB_STREAM_VERDICT`. The verdict can be `__SK_DROP`, `__SK_PASS`, or `__SK_REDIRECT`. Another use case for this program type is in the `strparser` framework (<https://www.kernel.org/doc/Documentation/networking/strparser.txt>). BPF programs can be used to parse message streams via callbacks for read operations, verdict and read completion. TLS and KCM use stream parsing.
- How do I attach my program? A redirection program is attached to a sockmap as `BPF_SK_SKB_STREAM_VERDICT`; it should return the result of `bpf_sk_redirect_map()`. A `strparser` program is attached via `BPF_SK_SKB_STREAM_PARSER` and should return the length of data parsed.
- What context is provided? A pointer to the struct `__sk_buff` containing packet metadata/data. However more fields are accessible to the `sk_skb` program type. The extra set of fields available are documented in `include/linux/bpf.h` like so:

```
/* Accessed by BPF_PROG_TYPE_sk_skb types from here to ... */
u32 family;
u32 remote ip4; /* Stored in network byte order */
u32 local ip4; /* Stored in network byte order */
u32 remote ip6[4]; /* Stored in network byte order */
u32 local ip6[4]; /* Stored in network byte order */
u32 remote port; /* Stored in network byte order */
u32 local port; /* stored in host byte order */
/* ... here. */
```

So from the above alone we can see we can gather information about the socket, since the above represents the key information that identifies the socket uniquely (protocol is already available in the globally-accessible portion of the struct `__sk_buff`).

- When does it run? A stream parser can be attached to a socket via `BPF_SK_SKB_STREAM_PARSER` attachment to a sockmap, and the parser runs on socket receive via `smap_parse_func_strparser()` in `kernel/bpf/sockmap.c`. `BPF_SK_SKB_STREAM_VERDICT` also attaches to the sockmap, and is run via `smap_verdict_func()`.